Nishant Balepur

Email: nbalepur@umd.edu, nishantbalepur@gmail.com
Website: nbalepur.github.io

RESEARCH SUMMARY

I am a computer science Ph.D. candidate at the University of Maryland, advised by Professors Jordan Boyd-Graber and Rachel Rudinger. I am interested in developing text generation and question answering systems that help users achieve their goals, with a focus on improving factuality, reasoning, preference training, alignment, and evaluations.

EDUCATION

- **University of Maryland, College Park (UMD)** — College Park, MD
  Ph.D. Computer Science; GPA: 4.00/4.00 — Aug 2023 - May 2027 (Expected)
  Advisors: Jordan Boyd-Graber, Rachel Rudinger
  Thesis (Proposed): Teaching AI to Answer Questions with Reasoning that Actually Helps You
  Committee: Jordan Boyd-Graber, Rachel Rudinger, Shi Feng, Fumeng Yang, David Weintrop

- **New York University** — New York City, NY
  Visiting Student — Aug 2025 - May 2026
  Host: Eunsol Choi

- **University of Illinois at Urbana-Champaign (UIUC)** — Urbana, IL
  B.S. Computer Science; B.S. Statistics (Dual Degree); GPA: 4.00/4.00 — Aug 2019 - May 2023
  Collaborators: Kevin Chen-Chuan Chang, Jiawei Han, Hari Sundaram, Diyi Yang

SELECTED WORKS

- Which of These Best Describes Multiple Choice Evaluations? A) Forced B) Flawed C) Fixable D) All of the Above
  Under Review (ACL), Best Paper Award at MASC-SLL 2025, ACL 2025 Best Paper Award Nomination
  Nishant Balepur, Rachel Rudinger, Jordan Boyd-Graber
  TL;DR: We critique multiple-choice evaluations with LLMs, and propose solutions from education to remedy them

- A Smart Mnemonic Sounds like Glue Tonic: Mixing LLMs with Student Feedback to Make Mnemonic Learning Stick
  EMNLP 2024
  Nishant Balepur, Matthew Shu, Alexander Hoyle, ..., Shi Feng, Seraphina Goldfarb-Tarrant, Jordan Boyd-Graber
  TL;DR: We use DPO to generate mnemonics aligned with what users prefer and aid learning

- Artifacts or Abduction: How Do LLMs Answer Multiple-Choice Questions Without the Question?
  ACL 2024, Best Paper Award at MASC-SLL 2024
  Nishant Balepur, Abhilasha Ravichander, Rachel Rudinger
  TL;DR: We discover LLMs score highly in MCQA benchmarks without the question, and analyze how

PUBLICATIONS

* denotes equal contribution, † denotes mentored student

1. A Good Plan is Hard to Find: Aligning LLMs with Preferences is Misaligned with Helpfulness
   Under Review (EMNLP)
   Nishant Balepur, Matthew Shu†, ..., Shi Feng, Fumeng Yang, Rachel Rudinger, Jordan Boyd-Graber

2. A Survey on LLM Jailbreaking Attacks and Defenses
   Under Review (COLM)
   Sander Schulhoff*, **Nishant Balepur***, Arjun Akkiraju*, Michael Ilie*, ..., Rob Voigt, Denis Peskoff

3. Which of These Best Describes Multiple Choice Evaluations? A) Forced B) Flawed C) Fixable D) All of the Above
   Under Review (ACL), Best Paper Award at MASC-SLL 2025, ACL 2025 Best Paper Award Nomination
   Nishant Balepur, Rachel Rudinger, Jordan Boyd-Graber

4. Whose Boat Does it Float? Improving Personalization in Preference Tuning via Inferred User Personas
   Under Review (ACL)
   Nishant Balepur, Vishakh Padmakumar, Fumeng Yang, Shi Feng, Rachel Rudinger, Jordan Boyd-Graber

5. Reverse Question Answering: Can an LLM Write a Question so Hard (or Bad) that it Can't Answer?
   NAACL 2025, Oral
   Nishant Balepur, Feng Gu, Abhilasha Ravichander, Shi Feng, Jordan Boyd-Graber, Rachel Rudinger

6. MoDS: Moderating a Mixture of Document Speakers to Summarize Debatable Queries in Document Collections
   NAACL 2025, Oral
   Nishant Balepur, Alexa Siu, Nedim Lipka, Franck Dernoncourt, Tong Sun, Jordan Boyd-Graber, Puneet Mathur

7. A Smart Mnemonic Sounds like Glue Tonic: Mixing LLMs with Student Feedback to Make Mnemonic Learning Stick
   EMNLP 2024, Media Coverage
   Nishant Balepur, Matthew Shu†, Alexander Hoyle, ..., Shi Feng, Seraphina Goldfarb-Tarrant, Jordan Boyd-Graber

8. KARL: Knowledge-Aware Retrieval and Representations aid Retention and Learning in Students
   EMNLP 2024
   Matthew Shu*†, **Nishant Balepur***, Shi Feng*, Jordan Boyd-Graber

9. Plausibly Problematic Questions in Multiple-Choice Benchmarks for Commonsense Reasoning
   EMNLP 2024 (Findings)
   Shramay Palta, **Nishant Balepur**, Peter Rankel, Sarah Wiegreffe, Marine Carpuat, Rachel Rudinger

10. The Prompt Report: A Systematic Survey of Prompting Techniques
    Preprint, Huggingface #1 Paper of the Day
    Sander Schulhoff*, Michael Ilie*, **Nishant Balepur**, ..., Shyamal Anadkat, Alexander Hoyle, Phillip Resnik

11. Artifacts or Abduction: How Do LLMs Answer Multiple-Choice Questions Without the Question?
    ACL 2024, Best Paper Award at MASC-SLL 2024
    Nishant Balepur, Abhilasha Ravichander, Rachel Rudinger

12. It's Not Easy Being Wrong: Large Language Models Struggle with Process of Elimination Reasoning
    ACL 2024 (Findings)
    Nishant Balepur, Shramay Palta, Rachel Rudinger

13. Is Your Large Language Model Knowledgeable or a Choices-Only Cheater?
    ACL 2024 (KnowLLM Workshop)
    Nishant Balepur, Rachel Rudinger

14. Expository Text Generation: Imitate, Retrieve, Paraphrase
    EMNLP 2023
    Nishant Balepur, Jie Huang, Kevin Chen-Chuan Chang

15. Text Fact Transfer
    EMNLP 2023
    Nishant Balepur, Jie Huang, Kevin Chen-Chuan Chang

16. DynaMiTE: Discovering Explosive Topic Evolutions with User Guidance
    ACL 2023 (Findings)
    Nishant Balepur*, Shivam Agarwal*, Karthik Ramanan, Susik Yoon, Diyi Yang, Jiawei Han

INDUSTRY RESEARCH EXPERIENCE

- **Allen Institute for AI (Ai2)**   Seattle, WA
  Research Scientist Intern: Semantic Scholar (Personalized QA)   May 2025 - Aug 2025

- **Adobe**   San Jose, CA
  Research Scientist Intern: Document Intelligence (Multi-LLM QA)   May 2024 - Aug 2024
  Outcome: Patent Submission, Paper at NAACL 2025

- **Meta**   Menlo Park, CA
  Software Engineering Intern: Facebook Creators Well-being (Comment Safety)   May 2022 - Aug 2022
  Outcome: 15% drop in negative interactions among Facebook's 1 billion users

INVITED TALKS

- **Can an LLM Write a Question so Hard (or Bad) that it Can't Answer?**   Albuquerque, NM
  NAACL 2025   May 2025

- **MoDS: Summarizing Debatable Queries in Document Collections**   Albuquerque, NM
  NAACL 2025   May 2025

- **Which of These Best Describes Multiple-Choice Evaluation with LLMs?**   State College, PA
  Penn State University   April 2025

- **Which of These Best Describes Multiple-Choice Evaluation with LLMs?**   Remote
  Google Translate Reading Group   March 2025

- **Aligning LLMs to Answer Questions with Reasoning That Actually Helps You**   London, U.K.
  Imperial College, London   March 2025

- **How Do LLMs Answer Multiple-Choice Questions Without the Question?**   Baltimore, MD
  Johns Hopkins University   May 2024

FELLOWSHIPS AND GRANTS

- **NSF Graduate Research Fellowship Program (GRFP)**   April 2023 - April 2028
  Wrote proposal on NLP for information accessibility—$159,000 over 3 Years of Ph.D.

- **Cohere for AI Research Grant Program**   April 2024
  Wrote proposal on LLMs for AI Safety + Education—full access to Cohere models

- **Dean's Fellowship**   April 2023 - April 2025
  Awarded the Dean's Fellowship from UMD for outstanding academic achievement

Awards

- **MASC-SLL 2025 Best Paper Award** — April 2025
  Best paper award for our MCQA Position Paper by PSU
- **MASC-SLL 2024 Best Paper Award** — May 2024
  Best paper award for "Artifacts or Abduction" by JHU
- **UIUC Computer Science Graduation with Highest Honors** — May 2023
  Recommended by the UIUC computer science department to graduate with highest honors
- **C.W. Gear Outstanding Undergraduate Student** — May 2022
  Awarded to two seniors that have demonstrated excellence in research and service

Student Mentees

- **Matthew Shu** (B.S. Yale → M.S. Yale), 2023-Present, LLMs for Education
  Two papers at EMNLP 2024 (long, main)

- **Atrey Desai** (B.S. UMD), 2024-Present, LLM Evaluation

- **Jerry He** (HS Student → B.S. GTech), 2024-2025, LLMs for Education
  Submitted Project on LLMs+Education for Marriott's Ridge High School Gifted and Talented Research Program

Professional Service

- **Conference Reviewer** — UMD, 2022-2024
  Reviewer for: *ACL/ARR 2023-Present, COLING 2024, IEEE TASLP 2024, TrustNLP 2024
  Great Reviewer Nomination: April, June, August 2024 ARR
  Outstanding Reviewer: EMNLP 2024
- **Conference Volunteer** — UMD, 2023, 2025
  Volunteer at ACL 2023, NAACL 2025
- **Visiting Student Day Volunteer** — UMD, Mar 2024
  Volunteer and ambassador for UMD's visiting student day
- **Winter Storm LLM Workshop** — UMD, Jan 2023
  Led a 5-day workshop on LLMs for non-CS graduate students
- **SIGNLL** — UIUC, Aug 2020 - May 2021
  President of Special Interest Group for Natural Language Learning

Skills

- **Languages:** Python, JavaScript/HTML/CSS R, C++, Java, OCaml

- **Libraries:** Huggingface, Datasets, TRL, Pytorch, nltk, Spacy, BeautifulSoup