

# Nishant Balepur

Emails: nbalepur@umd.edu, nishantbalepur@gmail.com

Website: nbalepur.github.io

## RESEARCH INTERESTS

---

I am a first-year Ph.D. student in computer science at the University of Maryland, College Park, working with Professors Jordan Boyd-Graber and Rachel Rudinger. I conduct research with the goal of **aligning LLMs with human values**, with a focus on **factuality** in text generation, **human-centered** frameworks, and LLM **safety** and **interpretability**. I am extremely grateful to be funded by the NSF Graduate Research Fellowship Program.

## EDUCATION

---

- **University of Maryland, College Park (UMD)** College Park, MD  
Ph.D. Computer Science Aug 2023 - Present  
Advisors: Professors Jordan Boyd-Graber, Rachel Rudinger
- **University of Illinois at Urbana-Champaign (UIUC)** Urbana, IL  
B.S. Computer Science; B.S. Statistics (Dual Degree); GPA: 4.00/4.00 Aug 2019 - May 2023  
Collaborators: Professors Kevin Chen-Chuan Chang, Jiawei Han, Hari Sundaram, Diyi Yang

## PUBLICATIONS AND WRITTEN WORK

---

- It's Not Easy Being Wrong: Evaluating Process of Elimination Reasoning in Large Language Models  
*arXiv:2311.07532*  
Nishant Balepur, Shramay Palta, Rachel Rudinger  
TL;DR: We uncover a new weakness of LLMs – reasoning toward incorrect options on multiple-choice questions
- Expository Text Generation: Imitate, Retrieve, Paraphrase  
*EMNLP 2023*  
Nishant Balepur, Jie Huang, Kevin Chen-Chuan Chang  
TL;DR: We design a task and model to generate multi-sentence and stylistically consistent factual texts
- Text Fact Transfer  
*EMNLP 2023*  
Nishant Balepur, Jie Huang, Kevin Chen-Chuan Chang  
TL;DR: We propose a complement to style transfer, where models must preserve style while transferring facts
- DynaMiTE: Discovering Explosive Topic Evolutions with User Guidance  
*Findings of ACL 2023*  
Nishant Balepur\*, Shivam Agarwal\*, Karthik Ramanan, Susik Yoon, Diyi Yang, Jiawei Han  
TL;DR: We build a model to mine for topics evolutions in large corpora, leveraging user-provided seed guidance
- Mastering the ABCDs of Complex Questions: Answer-Based Claim Decomposition for Self-Evaluating LLMs  
*arXiv:2305.14750*  
Nishant Balepur, Jie Huang, Samraj Moorjani, Kevin Chen-Chuan Chang, Hari Sundaram  
TL;DR: We study whether LLMs can perform a fine-grained form of self-evaluation
- Aligning Language Models with Factuality and Truthfulness  
*Undergraduate Senior Thesis*  
Nishant Balepur, Kevin Chen-Chuan Chang

## WORK EXPERIENCE

---

- **Meta** Menlo Park, CA  
Software Engineering Intern May 2022 - Aug 2022
- **HiMarley** Remote  
Data Science Intern May 2021 - Aug 2021
- **State Farm** Champaign, IL  
Actuarial and Modeling Intern Aug 2020 - Dec 2020
- **John Deere** Remote  
Software Engineering Intern Jun 2020 - Aug 2020

## STUDENTS MENTORED

---

- **Matthew Shu** (B.S. Yale), 2023-Present, NLP in Education  
Submitted a first-authored paper to NAACL 2024

## HONORS AND AWARDS

---

- **NSF Graduate Research Fellowship Program (GRFP)** April 2023 - April 2028  
*Provided 3 Years of Support over a 5 Year Period*
- **Dean's Fellowship** April 2023 - April 2025  
*Awarded the Dean's Fellowship from UMD for outstanding academic achievement*
- **UIUC Computer Science Graduation with Highest Honors** May 2023  
*Recommended by the UIUC computer science department to graduate with highest honors*
- **C.W. Gear Outstanding Undergraduate Student** May 2022  
*Awarded to two seniors that have demonstrated excellence in research and service*
- **James N. Snyder Memorial Award** May 2021  
*Awarded to three juniors based on academic merit*