

RESEARCH SUMMARY

I am a computer science Ph.D. candidate at the University of Maryland, advised by Professors Jordan Boyd-Graber and Rachel Rudinger. I broadly work on building text generation and question answering systems that help users achieve their goals, with a focus on improving factuality, reasoning, learning from human feedback, alignment, and evaluations.

EDUCATION

- University of Maryland, College Park (UMD)** College Park, MD
Ph.D. Computer Science; GPA: 4.00/4.00 Aug 2023 - May 2027 (Expected)
Advisors: Jordan Boyd-Graber, Rachel Rudinger
Thesis (Proposed): Teaching AI to Answer Questions with Reasoning that Actually Helps You
Committee: Jordan Boyd-Graber, Rachel Rudinger, Shi Feng, Fumeng Yang, David Weintrop
- New York University** New York City, NY
Visiting Student Aug 2025 - May 2026
Host: Eunsol Choi
- University of Illinois at Urbana-Champaign (UIUC)** Urbana, IL
B.S. Computer Science; B.S. Statistics (Dual Degree); GPA: 4.00/4.00 Aug 2019 - May 2023
Collaborators: Kevin Chen-Chuan Chang, Jiawei Han, Hari Sundaram, Diyi Yang

SELECTED WORKS

- Which of These Best Describes Multiple Choice Evaluations? A) Forced B) Flawed C) Fixable D) All of the Above
ACL 2025, **Best Paper Award at MASC-SLL 2025**, **ACL 2025 Best Paper Award Nomination**
Nishant Balepur, Rachel Rudinger, Jordan Boyd-Graber
TL;DR: We critique multiple-choice evaluations with LLMs, and propose solutions from education to remedy them
- A Smart Mnemonic Sounds like Glue Tonic: Mixing LLMs with Student Feedback to Make Mnemonic Learning Stick
EMNLP 2024
Nishant Balepur, Matthew Shu, Alexander Hoyle, ..., Shi Feng, Seraphina Goldfarb-Tarrant, Jordan Boyd-Graber
TL;DR: We use DPO to generate mnemonics aligned with what users prefer and aid learning
- Artifacts or Abduction: How Do LLMs Answer Multiple-Choice Questions Without the Question?
ACL 2024, **Best Paper Award at MASC-SLL 2024**
Nishant Balepur, Abhilasha Ravichander, Rachel Rudinger
TL;DR: We discover LLMs score highly in MCQA benchmarks without the question, and analyze how

CONFERENCE PAPERS

* denotes equal contribution, † denotes mentored student

- Which of These Best Describes Multiple Choice Evaluations? A) Forced B) Flawed C) Fixable D) All of the Above
ACL 2025, **Best Paper Award at MASC-SLL 2025**, **ACL 2025 Best Paper Award Nomination**
Nishant Balepur, Rachel Rudinger, Jordan Boyd-Graber
- Whose Boat Does it Float? Improving Personalization in Preference Tuning via Inferred User Personas
ACL 2025
Nishant Balepur, Vishakh Padmakumar, Fumeng Yang, Shi Feng, Rachel Rudinger, Jordan Boyd-Graber
- Reverse Question Answering: Can an LLM Write a Question so Hard (or Bad) that it Can't Answer?
NAACL 2025, **Oral**
Nishant Balepur, Feng Gu, Abhilasha Ravichander, Shi Feng, Jordan Boyd-Graber, Rachel Rudinger
- MoDS: Moderating a Mixture of Document Speakers to Summarize Debatable Queries in Document Collections
NAACL 2025, **Oral**
Nishant Balepur, Alexa Siu, Nedim Lipka, Franck Dernoncourt, Tong Sun, Jordan Boyd-Graber, Puneet Mathur
- A Smart Mnemonic Sounds like Glue Tonic: Mixing LLMs with Student Feedback to Make Mnemonic Learning Stick
EMNLP 2024, **Media Coverage**
Nishant Balepur, Matthew Shu†, Alexander Hoyle, ..., Shi Feng, Seraphina Goldfarb-Tarrant, Jordan Boyd-Graber
- KARL: Knowledge-Aware Retrieval and Representations aid Retention and Learning in Students
EMNLP 2024
Matthew Shu*†, **Nishant Balepur***, Shi Feng*, Jordan Boyd-Graber
- Plausibly Problematic Questions in Multiple-Choice Benchmarks for Commonsense Reasoning
EMNLP 2024 (Findings)
Shramay Palta, **Nishant Balepur**, Peter Rankel, Sarah Wiegrefe, Marine Carpuat, Rachel Rudinger

8. Artifacts or Abduction: How Do LLMs Answer Multiple-Choice Questions Without the Question?
ACL 2024, **Best Paper Award at MASC-SLL 2024**
Nishant Balepur, Abhilasha Ravichander, Rachel Rudinger
9. It's Not Easy Being Wrong: Large Language Models Struggle with Process of Elimination Reasoning
ACL 2024 (Findings)
Nishant Balepur, Shramay Palta, Rachel Rudinger
10. Expository Text Generation: Imitate, Retrieve, Paraphrase
EMNLP 2023, **Added to UIUC's AI Reading List for Qualifying Exams**
Nishant Balepur, Jie Huang, Kevin Chen-Chuan Chang
11. Text Fact Transfer
EMNLP 2023
Nishant Balepur, Jie Huang, Kevin Chen-Chuan Chang
12. DynaMiTE: Discovering Explosive Topic Evolutions with User Guidance
ACL 2023 (Findings)
Nishant Balepur*, Shivam Agarwal*, Karthik Ramanan, Susik Yoon, Diyi Yang, Jiawei Han

WORKSHOP PAPERS

* denotes equal contribution, † denotes mentored student

1. Language Models Generate Multiple-Choice Questions with Artifacts
MASC-SLL 2025
Atrey Desai†, **Nishant Balepur**, Rachel Rudinger
2. Mental Model for Machine Translation in Human-MT Scenario
MASC-SLL 2025
Hyojung Han, **Nishant Balepur**, Jordan Boyd-Graber, Marine Carpuat
3. Is Your Large Language Model Knowledgeable or a Choices-Only Cheater?
ACL 2024 (KnowLLM Workshop)
Nishant Balepur, Rachel Rudinger

PREPRINTS AND REPORTS

* denotes equal contribution, † denotes mentored student

1. Can They Dixit? Yes they Can! Dixit as a Playground for Multimodal Language Model Capabilities
Under Review (EMNLP Wordplay Workshop)
Nishant Balepur*, Dang Nguyen*, Dayeon Ki
2. A Good Plan is Hard to Find: Aligning Models with Preferences is Misaligned with What Helps Users
Under Review (EMNLP)
Nishant Balepur, Matthew Shu†, ..., Shi Feng, Fumeng Yang, Rachel Rudinger, Jordan Boyd-Graber
3. Mitigating Gender-Occupation Stereotypes in Llama through Consistency-Guided Finetuning
Under Review (EMNLP)
Haozhe An, **Nishant Balepur**, Christabel Acquaye, Rachel Rudinger
4. A Survey on LLM Jailbreaking Attacks and Defenses
Under Review (COLM)
Sander Schulhoff*, **Nishant Balepur***, Arjun Akkiraju*, Michael Ilie*, ..., Rob Voigt, Denis Peskoff
5. The Prompt Report: A Systematic Survey of Prompting Techniques
Preprint, **Huggingface #1 Paper of the Day**
Sander Schulhoff*, Michael Ilie*, **Nishant Balepur**, ..., Shyamal Anadkat, Alexander Hoyle, Phillip Resnik
6. Mastering the ABCDs of Complex Questions: Answer-Based Claim Decomposition for Self-Evaluating LLMs
Preprint
Nishant Balepur, Jie Huang, Samraj Moorjani, Kevin Chen-Chuan Chang, Hari Sundaram

INDUSTRY RESEARCH EXPERIENCE

- **Allen Institute for AI (Ai2)** Seattle, WA
Research Scientist Intern: Semantic Scholar (Personalized QA) May 2025 - Aug 2025
- **Adobe** San Jose, CA
Research Scientist Intern: Document Intelligence (Multi-LLM QA) May 2024 - Aug 2024
Outcome: Patent Submission, Paper at NAACL 2025
- **Meta** Menlo Park, CA
Software Engineering Intern: Facebook Creators Well-being (Comment Safety) May 2022 - Aug 2022
Outcome: 15% drop in negative interactions among Facebook's 1 billion users

STARTUP EXPERIENCE

- **LearnPrompting** Remote
Contractor: Co-leading research team on LLM jailbreak competitions
Designed harmfulness evaluations for HackAPrompt, lead author on two papers
Nov 2024 - Present
- **XpertNetwork** New York City
Consultant: Technical advice for startup using LLMs for expert network outreach
May 2025 - Present

INVITED TALKS

- **Can an LLM Write a Question so Hard (or Bad) that it Can't Answer?** Albuquerque, NM
NAACL 2025 May 2025
- **MoDS: Summarizing Debatable Queries in Document Collections** Albuquerque, NM
NAACL 2025 May 2025
- **Which of These Best Describes Multiple-Choice Evaluation with LLMs?** State College, PA
Penn State University April 2025
- **Which of These Best Describes Multiple-Choice Evaluation with LLMs?** Remote
Google Translate Reading Group March 2025
- **Aligning LLMs to Answer Questions with Reasoning That Actually Helps You** London, U.K.
Imperial College, London March 2025
- **How Do LLMs Answer Multiple-Choice Questions Without the Question?** Baltimore, MD
Johns Hopkins University May 2024

FELLOWSHIPS AND GRANTS

- **NSF Graduate Research Fellowship Program (GRFP)** April 2023 - April 2028
Wrote proposal on NLP for information accessibility—\$159,000 over 3 Years of Ph.D.
- **Cohere for AI Research Grant (2x)** April 2024
Wrote proposals on LLMs Education and HackAPrompt 2.0—full access to Cohere models
- **Dean's Fellowship** April 2023 - April 2025
Awarded the Dean's Fellowship from UMD for outstanding academic achievement

AWARDS

- **MASC-SLL 2025 Best Paper Award** April 2025
Best paper award for our MCQA Position Paper by PSU
- **MASC-SLL 2024 Best Paper Award** May 2024
Best paper award for “Artifacts or Abduction” by JHU
- **UIUC Computer Science Graduation with Highest Honors** May 2023
Recommended by the UIUC computer science department to graduate with highest honors
- **C.W. Gear Outstanding Undergraduate Student** May 2022
Awarded to two seniors that have demonstrated excellence in research and service

STUDENT MENTEES

- **Michael Xie** (M.S. UMD), 2025-Present, LLMs for Education
- **Matthew Shu** (B.S. Yale → M.S. Yale), 2023-Present, LLMs for Education
Two papers at EMNLP 2024 (long, main); one paper under review at EMNLP 2025 (long)
- **Atrey Desai** (B.S. UMD), 2024-Present, LLM Evaluation
One paper at MASC-SLL 2025 Workshop
- **Jerry He** (HS Student → B.S. GTech), 2024-2025, LLMs for Education
Submitted Project on LLMs+Education for Marriott's Ridge High School Gifted and Talented Research Program

PROFESSIONAL SERVICE

- **Conference Reviewer** UMD
Reviewer for: *ACL/ARR 2023-Present, IEEE TASLP 2024, TrustNLP 2024, SRW 2025
Great Reviewer Nomination: Apr, June (#6), Aug (#4), Oct (#2), Dec 2024, Feb 2025 ARR
Outstanding Reviewer: EMNLP 2024
2022-2024
- **Conference Volunteer** UMD
ACL 2023, NAACL 2025
2023, 2025
- **Visiting Student Day Volunteer** UMD
Volunteer and ambassador for UMD's visiting student day
Mar 2024, 2025

- **Winter Storm LLM Workshop**
Led a 5-day workshop on LLMs for non-CS graduate students
- **SIGNLL**
President of Special Interest Group for Natural Language Learning

UMD
Jan 2023

UIUC
Aug 2020 - May 2021

SKILLS

- **Languages:** Python, JavaScript/HTML/CSS R, C++, Java, OCaml
- **Libraries:** Huggingface, Datasets, TRL, Pytorch, nltk, Spacy, BeautifulSoup